

AI/LLM Intern Assessment Task

Task Title:

Build a Simple AI Q&A; Assistant using LLM API

Objective:

- 1 Assess Python fundamentals
- 2 Evaluate understanding of LLM APIs
- 3 Test basic RAG (Retrieval-Augmented Generation) knowledge
- 4 Understand problem-solving approach

Problem Statement:

Build a simple AI assistant that answers user questions based on a given document.

Requirements:

1. Input Data

- 1 Use any sample document (PDF, text, or webpage content)
- 2 Example: blog, company info, or dataset

2. Core Functionality

- 1 User inputs a question
- 2 System returns an answer using an LLM

3. LLM Integration (Mandatory)

- 1 Use OpenAI API, Gemini API, or Claude

4. Basic RAG Implementation

- 1 Split document into chunks
- 2 Generate embeddings
- 3 Retrieve relevant context
- 4 Pass context + query to LLM

5. Interface

- 1 CLI or simple UI (Streamlit optional)

6. Code Requirements

- 1 Clean and modular Python code
- 2 Proper function usage
- 3 Push code to GitHub

Bonus (Optional)

- 1 Use LangChain or LangGraph
- 2 Use vector databases like ChromaDB, FAISS, Pinecone
- 3 Add chat memory
- 4 Deploy the application

Submission Requirements

- 1 GitHub repository link
- 2 README with setup instructions
- 3 Sample queries and outputs

Evaluation Criteria

- 1 Python fundamentals
- 2 LLM workflow understanding
- 3 RAG implementation clarity
- 4 Code quality
- 5 Creativity and experimentation